

Redrob AI Candidate Ranking

Intelligent Discovery & Matching
System

What we built, why we built it, and
how it works.

What We Built & Why

- The Challenge:
 - Keyword matching is easily gamed and misses true context.
- The Philosophy (Why we built it this way):
 - Holistic View: Use all 23 behavioral signals (response rate, notice period, etc.).
 - Anti-Gaming: Trust-weighted skills and deep honeypot detection.
 - Context-Aware: Favor product-company experience and production deployment over consulting-heavy static roles.

How It Works: 9-Component Scorer

- A balanced, composite evaluation (Sum = 100%):
 - 25% - Title & Career (Role relevance, seniority, progression)
 - 18% - Skill Match (Proficiency $\times$ Duration $\times$ Endorsements)
 - 12% - Semantic Similarity (all-mpnet-base-v2 embeddings)
 - 10% - Experience Fit (Gaussian sweet spot around 5-9 yrs)
 - 10% - Behavioral Signals (Response rate, recency,

The Secret Sauce: Deep Analysis

- Going beyond standard filtering:
 - Job-Hopper Detection: Average tenure < 18 months incurs a scaling penalty.
 - Consulting vs. Product: Rewards product company backgrounds as requested by the JD.
 - Production Keywords: Scans descriptions for 'deployed', 'shipped', 'scale', etc.
 - Unique Fact-Based Reasoning: The LLM reasoning generator references actual skills, assessment scores, and raises honest concerns instead of using identical templates.

Security & Validity: 9 Honeypot Checks

- Detecting fakes and bot profiles:
 - Impossible YOE vs Education timeline
 - Excessive 'Expert' skill inflation for low YOE
 - Mass zero-duration expert skills
 - Reversed chronological timelines
 - Chronological impossibilities ($YOE > \text{max possible}$)
 - GitHub link integrity anomalies
 - Admin domain mismatch (Keyword stuffing)
 - Impossible behavioral bot patterns